

## ✓ Bank Marketing Dataset - Decision Tree Prediction

Performing prediction on a sample customer from the Bank Marketing dataset using a trained Decision Tree classifier, with label encoding for categorical features.

### ✓ 1 - Import libraries


```
# Core
import pandas as pd
import numpy as np

# Viz
import matplotlib.pyplot as plt
import seaborn as sns

# ML
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, LabelEncoder, StandardScaler
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, ConfusionMatrixDisplay
```

### ✓ 2 - Load the dataset

```
data = pd.read_csv('bank-full.csv', sep=';')
data.head()
```




	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	p
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	1
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	1
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	1
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	1
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	1

Next steps:

[Generate code with data](#)[View recommended plots](#)[New interactive sheet](#)


### ✓ 3 - Data Inspection & Cleaning

```
# Shape of dataset (rows, columns)
print("Dataset shape:", data.shape)
```



Dataset shape: (45211, 17)

```
# Preview first 5 rows
print("\nFirst 5 rows:")
display(data.head())
```



First 5 rows:

	age	job	marital	education	default	balance	housing	loan	contact	day	month	duration	campaign	pdays	previous	p
0	58	management	married	tertiary	no	2143	yes	no	unknown	5	may	261	1	-1	0	1
1	44	technician	single	secondary	no	29	yes	no	unknown	5	may	151	1	-1	0	1
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknown	5	may	76	1	-1	0	1
3	47	blue-collar	married	unknown	no	1506	yes	no	unknown	5	may	92	1	-1	0	1
4	33	unknown	single	unknown	no	1	no	no	unknown	5	may	198	1	-1	0	1

```
# Data types and non-null count
print("\nData Info:")
print(data.info())
```



```
Data Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays      45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
None
```

```
# Check for missing values
print("\nMissing values in each column:")
print(data.isnull().sum())
```



```
Missing values in each column:
age         0
job         0
marital     0
education   0
default     0
balance     0
housing     0
loan        0
contact     0
day         0
month       0
duration    0
campaign    0
pdays      0
previous    0
poutcome    0
y           0
dtype: int64
```

```
# Basic statistics for numeric columns
print("\nStatistical Summary (Numeric Columns):")
display(data.describe())
```



Statistical Summary (Numeric Columns):

	age	balance	day	duration	campaign	pdays	previous
<b>count</b>	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
<b>mean</b>	40.936210	1362.272058	15.806419	258.163080	2.763841	40.197828	0.580323
<b>std</b>	10.618762	3044.765829	8.322476	257.527812	3.098021	100.128746	2.303441
<b>min</b>	18.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000	0.000000
<b>25%</b>	33.000000	72.000000	8.000000	103.000000	1.000000	-1.000000	0.000000
<b>50%</b>	39.000000	448.000000	16.000000	180.000000	2.000000	-1.000000	0.000000
<b>75%</b>	48.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000	0.000000
<b>max</b>	95.000000	102127.000000	31.000000	4918.000000	63.000000	871.000000	275.000000

```
# Check unique values for each column
print("\nUnique values per column:")
for col in data.columns:
    print(f"{col}: {data[col].nunique()} unique values")
```



```
Unique values per column:
age: 77 unique values
job: 12 unique values
marital: 3 unique values
```

```

education: 4 unique values
default: 2 unique values
balance: 7168 unique values
housing: 2 unique values
loan: 2 unique values
contact: 3 unique values
day: 31 unique values
month: 12 unique values
duration: 1573 unique values
campaign: 48 unique values
pdays: 559 unique values
previous: 41 unique values
poutcome: 4 unique values
y: 2 unique values

```

## 4 - Preprocessing (Encoding + Splitting)

```

# Copy dataset to avoid modifying original
df = data.copy()

# Encode all categorical columns using LabelEncoder
label_encoders = {}
for col in df.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le

# Separate features (X) and target (y)
X = df.drop('y', axis=1) # 'y' is the target column
y = df['y']

# Train-test split (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)

↗ Training set shape: (36168, 16)
Testing set shape: (9043, 16)

```

## 5 - Model Training

```

# Initialize the Decision Tree model
dt_model = DecisionTreeClassifier(
    criterion='entropy', # use information gain
    max_depth=None,      # let the tree grow fully
    random_state=42
)

# Train the model
dt_model.fit(X_train, y_train)

print("Model training complete!")

↗ Model training complete!

```

## 6 - Train the Decision Tree Classifier

```

# Create the Decision Tree model
clf = DecisionTreeClassifier(random_state=42)

# Train the model
clf.fit(X_train, y_train)

↗
DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)

```

## 7 - Model Evaluation

After training, we evaluate the model on the test data (X\_test, y\_test). We'll calculate:

Accuracy — percentage of correct predictions.

Classification Report — precision, recall, F1-score.

Confusion Matrix — to visualize correct vs. incorrect predictions.

```
# Predictions on test data
y_pred = dt_model.predict(X_test)
```

```
# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
```

Accuracy: 0.8758

```
# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	7952
1	0.49	0.49	0.49	1091
accuracy			0.88	9043
macro avg	0.71	0.71	0.71	9043
weighted avg	0.88	0.88	0.88	9043

```
# Confusion matrix
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[7389 563]
 [ 560 531]]
```

## 8 - Visualizing the Decision Tree

```
!pip install graphviz
```

Requirement already satisfied: graphviz in /usr/local/lib/python3.11/dist-packages (0.21)

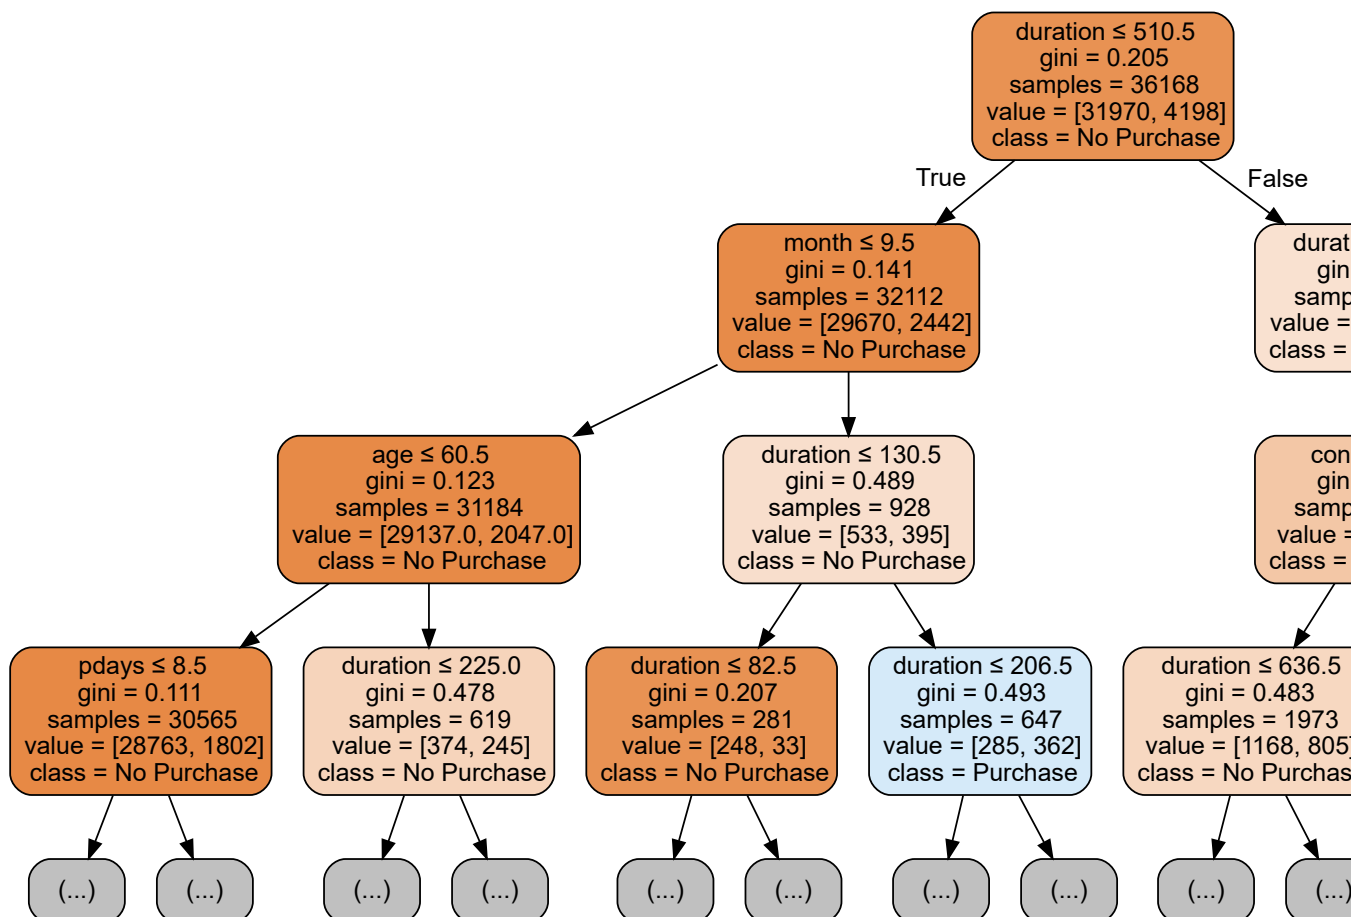
```
!apt-get install -y graphviz
```

Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done  
graphviz is already the newest version (2.42.2-6ubuntu0.1).  
0 upgraded, 0 newly installed, 0 to remove and 35 not upgraded.

```
from sklearn.tree import export_graphviz
import graphviz
```

```
dot_data = export_graphviz(
    clf,
    out_file=None,
    feature_names=X.columns,
    class_names=['No Purchase', 'Purchase'],
    filled=True,
    rounded=True,
    special_characters=True,
    max_depth=3, # Limit the depth of the tree for better visualization
)
```

```
graph = graphviz.Source(dot_data)
graph
```



## ✓ 9 - Confusion Matrix Visualization (Counts + Percentages)

```

# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

# Create percentage version
cm_percent = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis] * 100


# Create side-by-side plots
fig, ax = plt.subplots(1, 2, figsize=(12, 5))

# --- Left: Raw counts ---
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax[0],
            xticklabels=['No Purchase', 'Purchase'],
            yticklabels=['No Purchase', 'Purchase'])
ax[0].set_title("Confusion Matrix (Counts)")
ax[0].set_xlabel("Predicted")
ax[0].set_ylabel("Actual")

# --- Right: Percentages ---
sns.heatmap(cm_percent, annot=True, fmt=".2f", cmap="Blues", ax=ax[1],
            xticklabels=['No Purchase', 'Purchase'],
            yticklabels=['No Purchase', 'Purchase'])
ax[1].set_title("Confusion Matrix (%)")
ax[1].set_xlabel("Predicted")
ax[1].set_ylabel("Actual")

# Adjust layout & show
plt.tight_layout()
plt.show()

```



## 10 - Predicting a new result

```
# Define the sample customer data
new_data_dict = {
    'age': [36],
    'job': [label_encoders['job'].transform(['technician'])[0]],
    'marital': [label_encoders['marital'].transform(['married'])[0]],
    'education': [label_encoders['education'].transform(['secondary'])[0]],
    'default': [label_encoders['default'].transform(['no'])[0]],
    'balance': [368],
    'housing': [label_encoders['housing'].transform(['yes'])[0]],
    'loan': [label_encoders['loan'].transform(['yes'])[0]],
    'contact': [label_encoders['contact'].transform(['unknown'])[0]],
    'day': [6],
    'month': [label_encoders['month'].transform(['may'])[0]],
    'duration': [1597],
    'campaign': [2],
    'pdays': [-1],
    'previous': [0],
    'poutcome': [label_encoders['poutcome'].transform(['unknown'])[0]]
}

# Convert to DataFrame
new_data_df = pd.DataFrame(new_data_dict)

# Predict using the trained Decision Tree model
prediction = dt_model.predict(new_data_df)[0]
probability = dt_model.predict_proba(new_data_df)[0][1] # Probability of positive class

# Decode the prediction back to original class names
predicted_class = label_encoders['y'].inverse_transform([prediction])[0]

# Display results
print(f"Prediction: {predicted_class}")
print(f"Probability of Purchase: {probability:.2%}")

Prediction: yes
Probability of Purchase: 100.00%
```